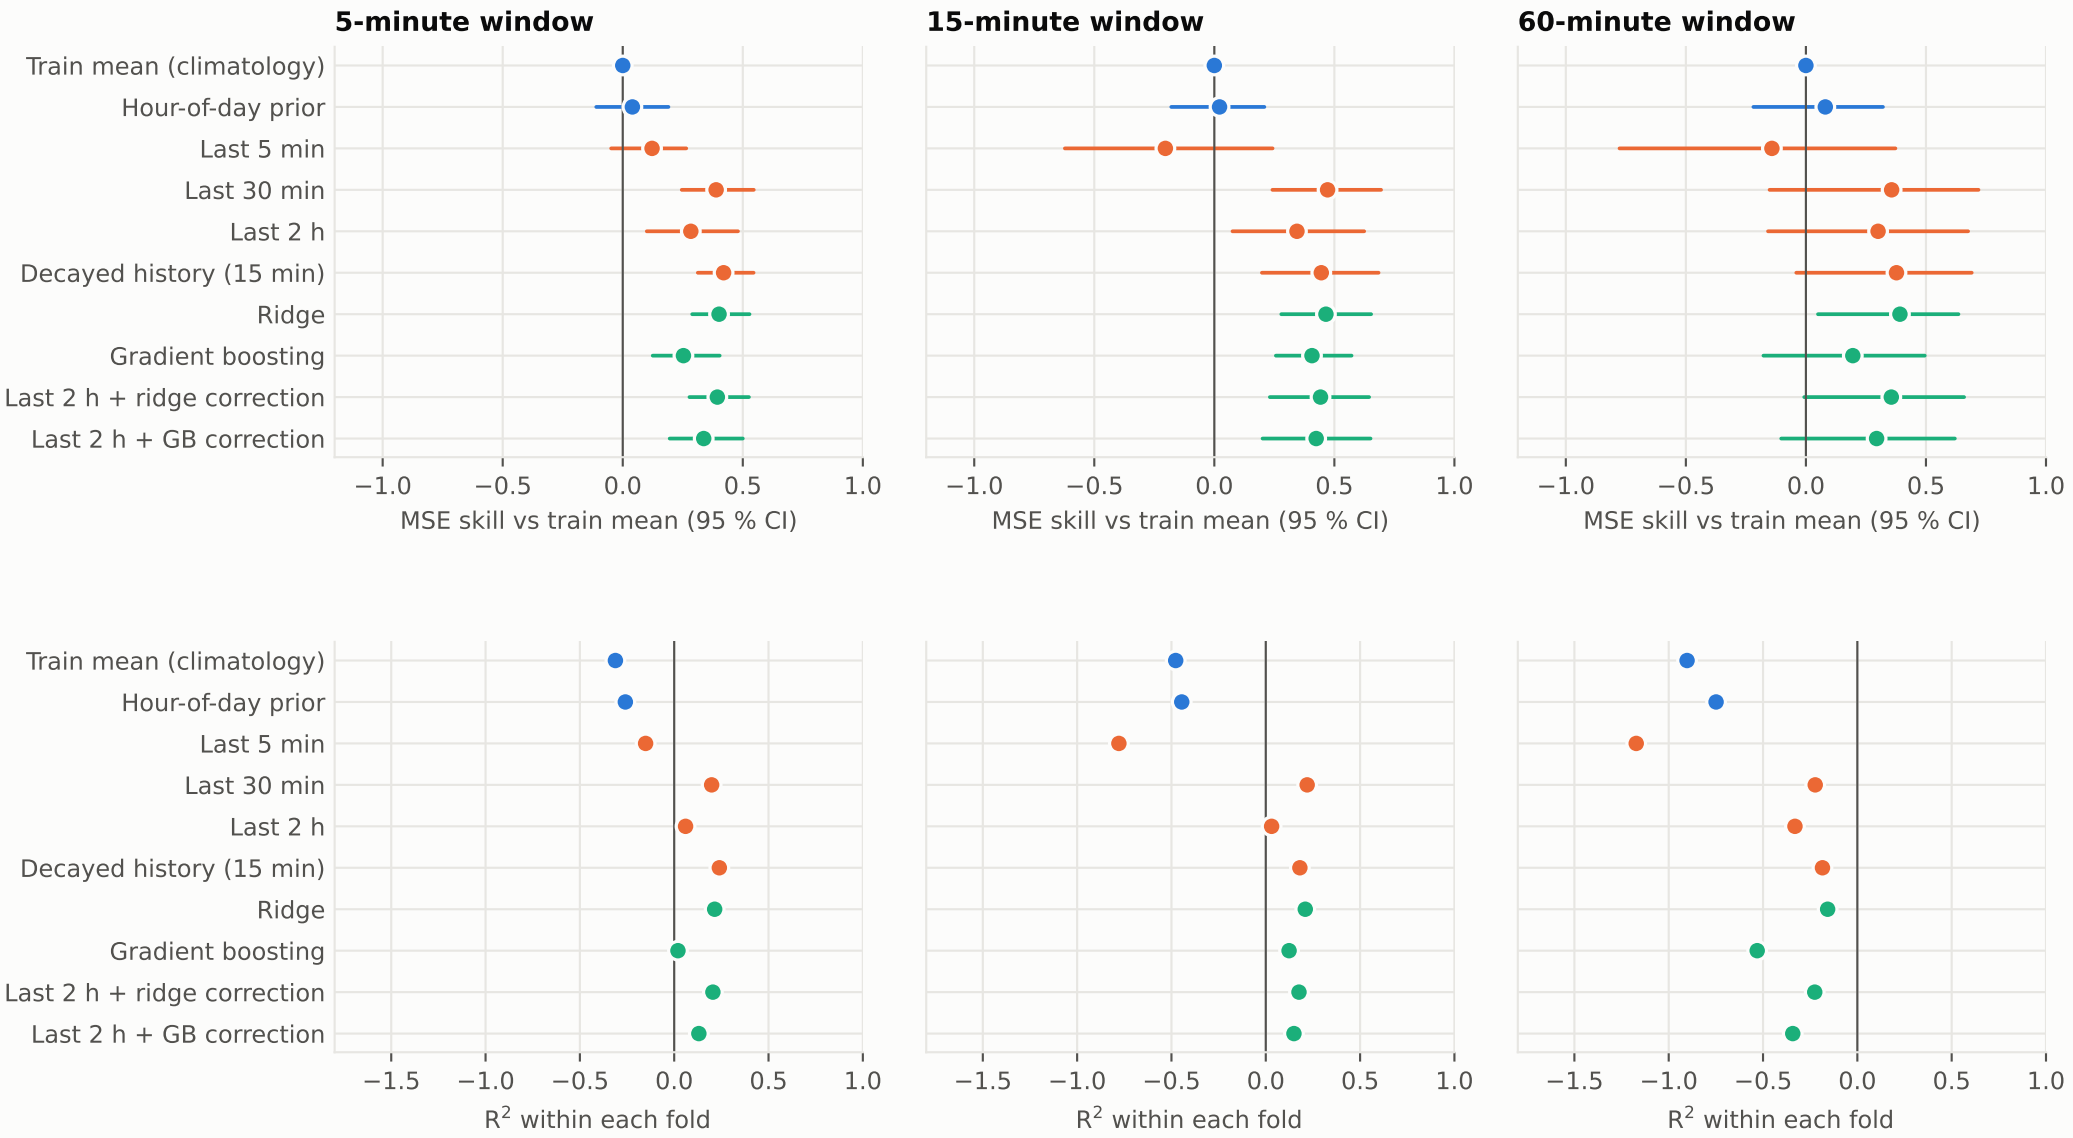
Second node LW1 (rural, 3 rolling folds, design frozen on CC1): method comparison, channel-6 band
5-minute window
15-minute window
60-minute window
Train mean (climatology)
Hour-of-day prior
Last 5 min
Last 30 min
Last 2 h
Decayed history (15 min)
Ridge
Gradient boosting
Last 2 h + ridge correction
Last 2 h + GB correction
−1.0
−0.5
0.0
0.5
1.0
MSE skill vs train mean (95 % CI)
−1.5
R² within each fold
Confirmation run, no tuning; 30-min and decayed memories score better than the 2 h average on point estimates, and learned models match but do not beat them; see docs/lw1_results.md.
Priors
History averages
Learned models